

CSE 100: Algorithm Design and Analysis

Chapter 24: Single-Source Shortest Paths

Sungjin Im

University of California, Merced

Last Update: 4-21-2023

Single-Source Shortest Paths

Problem definition

✓ Single-pair shortest-path problem:

Input: Directed graph $G = (V, E)$ with weight/distance $w(u, v)$ on each edge $(u, v) \in E$, and a pair of vertices $s, t \in V$.

Output: A shortest path from s to t .

Single-destination shortest-paths problem:

Input: G and w as above, and a destination vertex t .


Output: A shortest path to t from each vertex v .

✓ Single-source shortest-paths problem:

Input: G and w as above, and a source vertex s .

Output: A shortest path from s to each vertex v .

Preliminaries



Terminology:

- ▶ Path: a sequence of edges that connect a sequence of vertices.
So a path P can be represented as (v_0, v_1, \dots, v_k) where $v_0, v_1, \dots, v_k \in V$ and $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k) \in E$.
- ▶ Simple Path: a path P is said to be simple if no vertex appears more than once on the path.
 - * It is a convention that a path refers to a simple path; however, this is not necessarily the case in the textbook.
- ▶ The weight/distance of path P , $w(P)$ is defined as the total weight/distance of edges of the path P :
 $w(P) := w(v_0, v_1) + w(v_1, v_2) + \dots + w(v_{k-1}, v_k)$

Preliminaries

$u \rightsquigarrow v$: v is reachable from u . (sometimes, a path from u to v)


$u \rightsquigarrow^P v$: P is a path from u to v .

Shortest-path weight $\delta(u, v)$ from u to v is defined as

$$\delta(u, v) = \begin{cases} \min\{w(P) : u \rightsquigarrow^P v\}, & \text{if there is a path from } u \text{ to } v \\ \infty & \text{otherwise} \end{cases}$$

A shortest path from u to v is any path P from u to v such that $w(P)$ = $\delta(u, v)$.

Key Lemma



Lemma (24.1. Optimality of Subpaths)

If $P = \langle v_0, v_1, \dots, v_k \rangle$ is a shortest path from v_0 to v_k , then for any $0 \leq i \leq j \leq k$, $\langle v_i, v_{i+1}, \dots, v_j \rangle$ is a shortest path from v_i to v_j . In other words, subpaths of shortest paths are also shortest paths.



Proof.

Cut-and-paste. Otherwise, one could get a 'better' shortest path from v_0 to v_k by replacing $\langle v_i, v_{i+1}, \dots, v_j \rangle$ with a better path from v_i to v_j . □

Some Issues

What is the shortest distance from s to each vertex?

∞ : unreachable vertex for s .



✓ If there is a path from s to v including a negative-weight cycle, $f(s, v)$ is not well defined.

Some Issues



If there is a path from s to a vertex u that includes a negative-weight cycle, then a shortest path from s to v is not well-defined, and we have $\delta(s, v) = -\infty$.

Lemma

If there is a shortest path from u to v , then there is such a path that is simple.

Proof.

If the path contains a negative-weight cycle:

If the path contains a positive-weight cycle:

If the path contains a 0-weight cycle:



We can assume without loss of generality that a shortest path contains no cycles.

Some Issues

-

Corollary

If there is a shortest path from u to v , then there is such a path that is simple, therefore consists of at most $|V| - 1$ edges.

Representing Shortest paths from a Single Source s

A shortest path can be encoded by π : $v.\pi$ means v 's predecessor.
In path $\langle v_0 = s, v_1, v_2, \dots, v_k \rangle$, $v_0.\pi = \text{NIL}$, $v_1.\pi = v_0$, ..., $v_k.\pi = v_{k-1}$.

Definition

A shortest-paths tree rooted at s is a directed subgraph $G' = (V', E')$ where $V' \subseteq V$ and $E' \subseteq E$ such that

1. V' is the set of vertices reachable from s in G .
2. G' forms a rooted tree with root s , and
3. for all $v \in V'$, the unique simple path from s to $v \in G'$ is a shortest path from s to v in G .

Representing Shortest paths from a Single Source s

Lemma

Let $G = (V, E)$ be a weighted, directed graph with no negative-weight cycles reachable from source vertex $s \in V$. Then, there exists a shortest-paths tree rooted at s (over all reachable vertices from s).

Representing Shortest paths from a Single Source s

Lemma

Let $G = (V, E)$ be a weighted, directed graph with no negative-weight cycles reachable from source vertex $s \in V$. Then, there exists a shortest-paths tree rooted at s (over all reachable vertices from s).

Proof.

Use the optimality of subpaths lemma.



* Not necessarily unique.

Representing Shortest paths from a Single Source s

To compute the shortest path from source s to each vertex v (along with its distance), we only need to compute $v.d$ and $v.\pi$.
* $v.d = \infty$ implies that v is not reachable from s .

Key Subroutine: edge relaxation

First, initialize:

INITIALIZE-SINGLE-SOURCE(G, s)

1 **for** each vertex $v \in G.V$

2 $v.d = \infty$

3 $v.\pi = \text{NIL}$

4 $s.d = 0$

Think of $v.d$ as the shortest path (distance) estimate to v from s .

Key Subroutine: ~~edge relaxation~~



Relaxing edge (u, v) : if we can improve the (current) shortest path to v by replacing it with the (current) shortest path to u followed by (u, v) , then update $v.d$ and $v.\pi$ accordingly.

RELAX (u, v, w)

- 1 **if** $v.d > \underline{u.d + w(u, v)}$
- 2 $\underline{v.d} = u.d + w(u, v)$
- 3 $\underline{v.\pi} = u$

Key Subroutine: edge relaxation

Relaxing edge (u, v) : if we can improve the (current) shortest path to v by replacing it with the (current) shortest path to u followed by (u, v) , then update $v.d$ and $v.\pi$ accordingly.



Key Subroutine: edge relaxation

Most algorithms are based on relaxing edges (after the initialization).

- ▶ Dijkstra's algorithm and the shortest-paths algorithm for DAGs relax each edge exactly once.
- ▶ Bellman-Ford algorithm relaxes each edge $|V| - 1$ times.

Useful Properties

Lemma (24.10. Triangle Inequality)

For any edge (u, v) , $\delta(s, v) \leq \delta(s, u) + w(u, v)$.

Suppose we first set $v.d = \infty$ for all $v \in V$ except the source vertex, and update $v.d$ only via relaxing some edges. Then we have the following properties:

Lemma (24.11. Upper-bound Property)

We always have $v.d \geq \delta(s, v)$ for all $v \in V$, and once $v.d$ achieves value $\delta(s, v)$, it never changes.

Corollary (24.12. No-path Property)

If there is no path from s to v , then we always have $v.d = \delta(s, v) = \infty$.

Useful Properties



Lemma (24.14. Convergence Property)

If $s \rightsquigarrow u \rightarrow v$ is a shortest path in G for some $u, v \in G$, and if $u.d = \delta(s, u)$ at any time prior to relaxing edge (u, v) , then $v.d = \delta(s, v)$ at all times afterward.

Lemma (24.15. Path-relaxation Property)

If $P = \langle \underline{v_0}, v_1, \dots, \underline{v_k} \rangle$ is a shortest path from $s = v_0$ to v_k , and the sequence of relaxations includes $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$ as a subsequence, then $v_k.d = \delta(s, v_k)$.


Lemma (24.17. Predecessor-subgraph Property)

Once $v.d = \delta(s, v)$ for all $v \in V$, the predecessor subgraph is a shortest-paths tree rooted at s .

The Bellman-Ford Algorithm

- ▶ Solves the single-source shortest-paths problem in the general case in which edges can have negative weights.
- ▶ Returns false if there is a negative-weight cycle reachable from s . Otherwise returns true along with the shortest paths and their distances.

The Bellman-Ford Algorithm



BELLMAN-FORD(G, w, s)

1 INITIALIZE-SINGLE-SOURCE(G, s)

$V-1$ rounds

2 for $i = 1$ to $|G.V| - 1$

3 for each edge $(u, v) \in G.E$

4 RELAX(u, v, w)

5 for each edge $(u, v) \in G.E$

6 if $v.d > u.d + w(u, v)$

7 return FALSE

8 return TRUE

False

for \forall edge (u, v)

relax(u, v, w)

If any $u.d$ has
changed, return false.

\Rightarrow when G has a negative
weight cycle reachable
from s .


The Bellman-Ford Algorithm

Example

At any point in time,

$v.d = D$ means there's a path of distance D from s to v .

if $v.d = \infty$ means we haven't found a path from s to v .




The Bellman-Ford Algorithm

Example

At any point in time,

$v.d = D$ means there's a path of distance D from s to v .

if $v.d = \infty$ means we haven't found a path from s to v .



The Bellman-Ford Algorithm

Correctness: True case

At any point in time,

$v.d = D$ means there's a path of distance D from s to v .

if $v.d = \infty$ means we haven't found a path from s to v .

As claimed before all algorithms in this chapter are based on edge relaxation:

Lemma (24.11. Upper-bound Property)

We always have $v.d \geq \delta(s, v)$ for all $v \in V$, and once $v.d$ achieves value $\delta(s, v)$, it never changes.

Corollary (24.12. No-path Property)

If there is no path from s to v , then we always have $v.d = \delta(s, v) = \infty$.

The Bellman-Ford Algorithm

Correctness: True case

For any shortest (simple) path $P = \langle v_0 = s, v_1, v_2, \dots, v_k = u \rangle$ from s to u , there is a relaxation subsequence of $(v_0, v_1), (v_1, v_2), (v_2, v_3), \dots, (v_{k-1}, v_k)$ generated by the BF.

The Bellman-Ford Algorithm

Correctness: True case

For any shortest (simple) path $P = \langle v_0 = s, v_1, v_2, \dots, v_k = u \rangle$ from s to u , there is a relaxation subsequence of $(v_0, v_1), (v_1, v_2), (v_2, v_3), \dots, (v_{k-1}, v_k)$ generated by the BF. Note that $k \leq |V| - 1$.

The Bellman-Ford Algorithm

Correctness: True case

For any shortest (simple) path $P = \langle v_0 = s, v_1, v_2, \dots, v_k = u \rangle$ from s to u , there is a relaxation subsequence of $(v_0, v_1), (v_1, v_2), (v_2, v_3), \dots, (v_{k-1}, v_k)$ generated by the BF. Note that $k \leq |V| - 1$. Do you see why?

The Bellman-Ford Algorithm

Correctness: True case

For any shortest (simple) path $P = \langle v_0 = s, v_1, v_2, \dots, v_k = u \rangle$ from s to u , there is a relaxation subsequence of $(v_0, v_1), (v_1, v_2), (v_2, v_3), \dots, (v_{k-1}, v_k)$ generated by the BF. Note that $k \leq |V| - 1$. Do you see why?

Lemma (24.15. Path-relaxation Property)

If $P = \langle v_0, v_1, \dots, v_k \rangle$ is a shortest path from $s = v_0$ to v_k , and the sequence of relaxations includes $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$ as a subsequence, then $v_k.d = \delta(s, v_k)$.

Thanks to the observation and Lemma 24.15, we have $u.d = \delta(s, u)$ for all $u \in V$ at the end.

The Bellman-Ford Algorithm

Correctness: True case

We have $u.d = \delta(s, u)$ for all $u \in V$ at the end. So,
 $v.d \leq u.d + w(u, v)$ for all (u, v) by Triangle Inequality. So the BF
returns True.

The Bellman-Ford Algorithm

Correctness: False case

For the sake of contradiction, suppose the BF returns True for the False case.

Say $c = \langle v_0, v_1, \dots, v_k = v_0 \rangle$ be a negative-weight cycle (reachable from s), so we have,

$$w(v_0, v_1) + w(v_1, v_2) + \dots + w(v_k, v_0) < 0.$$

But we know

$$v_1.d \leq v_0.d + w(v_0, v_1)$$

$$v_2.d \leq v_1.d + w(v_1, v_2)$$

...

$$v_k.d \leq v_{k-1}.d + w(v_{k-1}, v_k)$$

$$v_0.d \leq v_k.d + w(v_k, v_0)$$

So, we have $0 \leq w(v_0, v_1) + w(v_1, v_2) + \dots + w(v_k, v_0)$, a contradiction.

The Bellman-Ford Algorithm

Running Time

BELLMAN-FORD(G, w, s)

1 INITIALIZE-SINGLE-SOURCE(G, s)

2 **for** $i = 1$ **to** $|G.V| - 1$

3 **for** each edge $(u, v) \in G.E$

4 RELAX(u, v, w)

5 **for** each edge $(u, v) \in G.E$

6 **if** $v.d > u.d + w(u, v)$

7 **return** FALSE

8 **return** TRUE



$\left. \begin{array}{l} \text{for each edge } (u, v) \in G.E \\ \text{RELAX}(u, v, w) \end{array} \right\} E \cdot V$

$\left. \begin{array}{l} \text{for each edge } (u, v) \in G.E \\ \text{if } v.d > u.d + w(u, v) \end{array} \right\} E$

$E \cdot V$

The Bellman-Ford Algorithm

Running Time

```
BELLMAN-FORD( $G, w, s$ )  
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )  
2  for  $i = 1$  to  $|G.V| - 1$   
3      for each edge  $(u, v) \in G.E$   
4          RELAX( $u, v, w$ )  
5  for each edge  $(u, v) \in G.E$   
6      if  $v.d > u.d + w(u, v)$   
7          return FALSE  
8  return TRUE
```

$O(EV)$

The Bellman-Ford Algorithm

Questions

Q. Suppose that there is an integer $m > 0$ such that there is a shortest path from s to each vertex v consisting of at most m edges. Further, we know the value of m . How many times do you need to iterate Lines 3 and 4?

```
BELLMAN-FORD( $G, w, s$ )  
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )  
2  for  $i = 1$  to  $|G.V| - 1$   
3      for each edge  $(u, v) \in G.E$   
4          RELAX( $u, v, w$ )  
5  for each edge  $(u, v) \in G.E$   
6      if  $v.d > u.d + w(u, v)$   
7          return FALSE  
8  return TRUE
```


The Bellman-Ford Algorithm

Questions

Q. We would like to find a shortest path from the single source to each vertex. If we see no change of $v.d$ for any vertex v , we can stop. True or False?

```
BELLMAN-FORD( $G, w, s$ )  
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )  
2  for  $i = 1$  to  $|G.V| - 1$   
3      for each edge  $(u, v) \in G.E$   
4          RELAX( $u, v, w$ )  
5  for each edge  $(u, v) \in G.E$   
6      if  $v.d > u.d + w(u, v)$   
7          return FALSE  
8  return TRUE
```

Single-source Shortest Paths in DAGs

Questions



A DAG has no cycle. So no worries about negative-weight cycles.

DAG-SHORTEST-PATHS(G, w, s)

- 1 topologically sort the vertices of G
- 2 **INITIALIZE-SINGLE-SOURCE**(G, s)
- 3 **for** each vertex u , taken in topologically sorted order
- 4 **for** each vertex $v \in G.Adj[u]$
- 5 **RELAX**(u, v, w)


Single-source Shortest Paths in DAGs

Example



Single-source Shortest Paths in DAGs

Example



Single-source Shortest Paths in DAGs

Correctness

Say $P = \langle v_0 = s, v_1, v_2, \dots, v_k = v \rangle$ is a shortest path from s to v . Then v_0, v_1, \dots, v_k must appear in the topological ordering in this order.

The relaxation sequence includes $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$ as a subsequence.

Hence if a vertex v is reachable from s , then $v.d = \delta(s, v)$ by the path-relaxation property.

Single-source Shortest Paths in DAGs

Running time

DAG-SHORTEST-PATHS(G, w, s)

- 1 topologically sort the vertices of G
- 2 INITIALIZE-SINGLE-SOURCE(G, s)
- 3 **for** each vertex u , taken in topologically sorted order
- 4 **for** each vertex $v \in G.Adj[u]$
- 5 RELAX(u, v, w)

Single-source Shortest Paths in DAGs

Running time

DAG-SHORTEST-PATHS(G, w, s)

```
1  topologically sort the vertices of  $G$ 
2  INITIALIZE-SINGLE-SOURCE( $G, s$ )
3  for each vertex  $u$ , taken in topologically sorted order
4      for each vertex  $v \in G.Adj[u]$ 
5          RELAX( $u, v, w$ )
```

$E + V$
 V
 E

Topological sort: $O(E + V)$.

Each edge is relaxed exactly once. So, $O(E)$ for all relaxations.

Thus, $O(E + V)$.

Single-source Shortest Paths when Edges Have Non-negative Weights

No worries about negative weight cycles since we have no negative-weight edges.

The Dijkstra algorithm

- ▶ maintains a set S of vertices whose shortest distances have been determined.
- ▶ grows S by adding a vertex $u \in V - S$ with the shortest distance estimate, $u.d$ (and relaxing all edges leaving u).
- ▶ is similar to Prim's algorithm.


Dijkstra Algorithm's Pseudocode

DIJKSTRA(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = G.V$ 
4  while  $Q \neq \emptyset$ 
5       $u = \text{EXTRACT-MIN}(Q)$ 
6       $S = S \cup \{u\}$ 
7      for each vertex  $v \in G.Adj[u]$ 
8          RELAX( $u, v, w$ )
```


The Dijkstra algorithm: conceptual illustration

Example



The Dijkstra algorithm: actual illustration

Example



Dijkstra Algorithm's correctness

Proof sketch:

Want to show: when u is extracted from Q , we have $u.d = \delta(s, u)$, assuming that S consists of vertices whose shortest distances have been determined.

If $u = s$ or u is unreachable from s :

Otherwise, what happens if $u.d \neq \delta(s, u)$?

Where do we use the fact that edge weights are non-negative?

Dijkstra Algorithm's running time

Running Time:

Dijkstra Algorithm performs $O(V)$ Extract-Min and $O(E)$ Decrease-key.

DIJKSTRA(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = G.V$ 
4  while  $Q \neq \emptyset$ 
5       $u = \text{EXTRACT-MIN}(Q)$ 
6       $S = S \cup \{u\}$ 
7      for each vertex  $v \in G.Adj[u]$ 
8          RELAX( $u, v, w$ )
```

Dijkstra Algorithm's running time

Running Time:

Dijkstra Algorithm performs $|V|$ Extract-Min and $|E|$ Decrease-key.

If min-priority queue is implemented by array,

- ▶ Extract-Min: $O(V)$.
- ▶ Decrease-Key: $O(1)$.

$$O(V^2 + E) = O(V^2).$$

Dijkstra Algorithm's running time

Running Time:

Dijkstra Algorithm performs $|V|$ Extract-Min and $|E|$ Decrease-key.

If min-priority queue is implemented by binary heap,

- ▶ Extract-Min: $O(\log V)$.
- ▶ Decrease-Key: $O(\log V)$.

$$O(V \log V + E \log V) = O(E \log V).$$

Dijkstra Algorithm's running time

Running Time:

Dijkstra Algorithm performs $|V|$ Extract-Min and $|E|$ Decrease-key.

If min-priority queue is implemented by Fibonacci heap,

- ▶ Extract-Min: $O(\log V)$.
- ▶ Decrease-Key: $O(1)$ (in an amortized sense).

$O(V \log V + E)$.